

AI - FOUNDATION AND APPLICATION

Instructor:


Assoc. Prof. Dr. Truong Ngoc Son

Chapter 2

Back Propagation

Outline

binary classification




- ❑ Multiclass Classification With Softmax regression
- ❑ Lost function – cross entropy
- ❑ Stochastic gradient descent – batch and mini-batch gradient descent
- ❑ Translating math into code




Multiclass Classification With Softmax regression

10 e'p



Multiclass classification example



Predictive output, o **Desired output, y**

0.9	1
0.7	0
0.5	0

N samples, K outputs

Cost / Loss

$$L = \frac{1}{N} \sum_{t=1}^N \sum_{k=1}^K (y_k^t - o_k^t)^2$$


sample class MSE

Multiclass classification with logistic regression

Sigmoid function is used for the neurons at the output layer

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ❑ It is ideally for two-class classification
- ❑ The outputs are independent
- ❑ The sigmoid may produces high probability for all classes, some of them, or none of them



0.9	0.9	0.3
0.8	0.2	0.2
0.6	0.5	0.1

0.9

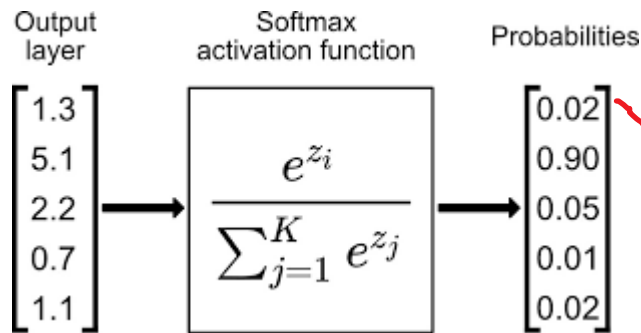
$$0.9 + 0.8 + 0.6$$

Multiclass classification with softmax regression

We expect that there is only one right answer, the outputs are mutually exclusive.

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

- ❑ The softmax will enforce that the sum of the probabilities of output classes are equal to one
- ❑ Softmax is used for multi-classification in the Logistic Regression model, whereas Sigmoid is used for binary classification in the Logistic Regression model



Handwritten red notes illustrating the calculation of the first probability:

$$e^{1.3}$$


$$e^{1.3} + e^{5.1} + e^{2.2} + e^{0.7} + e^{1.1}$$

Multiclass classification with softmax regression

- ❑ Softmax function is mostly used in a final layer of Neural Network
- ❑ The outputs are probability distribution

sigmoid: hidden layer - làm ngõ ra độc lập vs nhau

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$



Lost function – cross entropy

Cross-entropy takes the negative log likelihood of the predicted probability

Cross-entropy loss

$$Loss = - \sum_{i=1}^M y_i \log(o_i)$$


M – number of classes

y: class label

o: predicted probability observation


$$P(y_i = 1/x) = \frac{e^{w^{(i)T}x + b^{(i)}}}{\sum_{j=1}^k e^{w^{(j)T}x + b^{(j)}}}$$

$$C = \sum_{i=1}^k -y_i \log P(y_i = 1/x)$$



Back propagation

Feedforward propagation with softmax activation



$$Loss = - \sum_{k=1}^K y_k \log(o_k)$$


$$z_j^{(1)} = \sum_{i=1}^M x_i w_{j,i}^{(1)} + b_j^{(1)}$$

$$a_j^{(1)} = \frac{1}{1 + e^{-z_j^{(1)}}}$$

$$z_k^{(2)} = \sum_{j=1}^K a_j^{(1)} w_{k,j}^{(2)} + b_k^{(2)}$$

$$o_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j^{(2)}}}$$

Feedforward propagation with softmax activation



Gradient descent

$$w_{k,j} = w_{k,j} - \eta \frac{\partial L}{\partial w_{k,j}}$$

$$L(y, o) = - \sum_{k=1}^K y_k \log(o_k)$$

Apply the chain rule

$$\frac{\partial L}{\partial w_{k,j}} = \frac{\partial L}{\partial z_k} \frac{\partial z_k}{\partial w_{k,j}}$$

$$\frac{\partial L}{\partial z_k} = \frac{\partial L}{\partial o_k} \frac{\partial o_k}{\partial z_k}$$

$$o_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j^{(2)}}}$$

Error of k^{th} neuron of the output layer


$$\delta_k = \frac{\partial L}{\partial z_k}$$

Using calculus, we obtain $\frac{\partial L}{\partial z_k} = o_k - y_k$ $\delta_k = o_k - y_k$

o_k is the k th component of the neuron's prediction
and y_k is the k th component of the label

Back-propagation

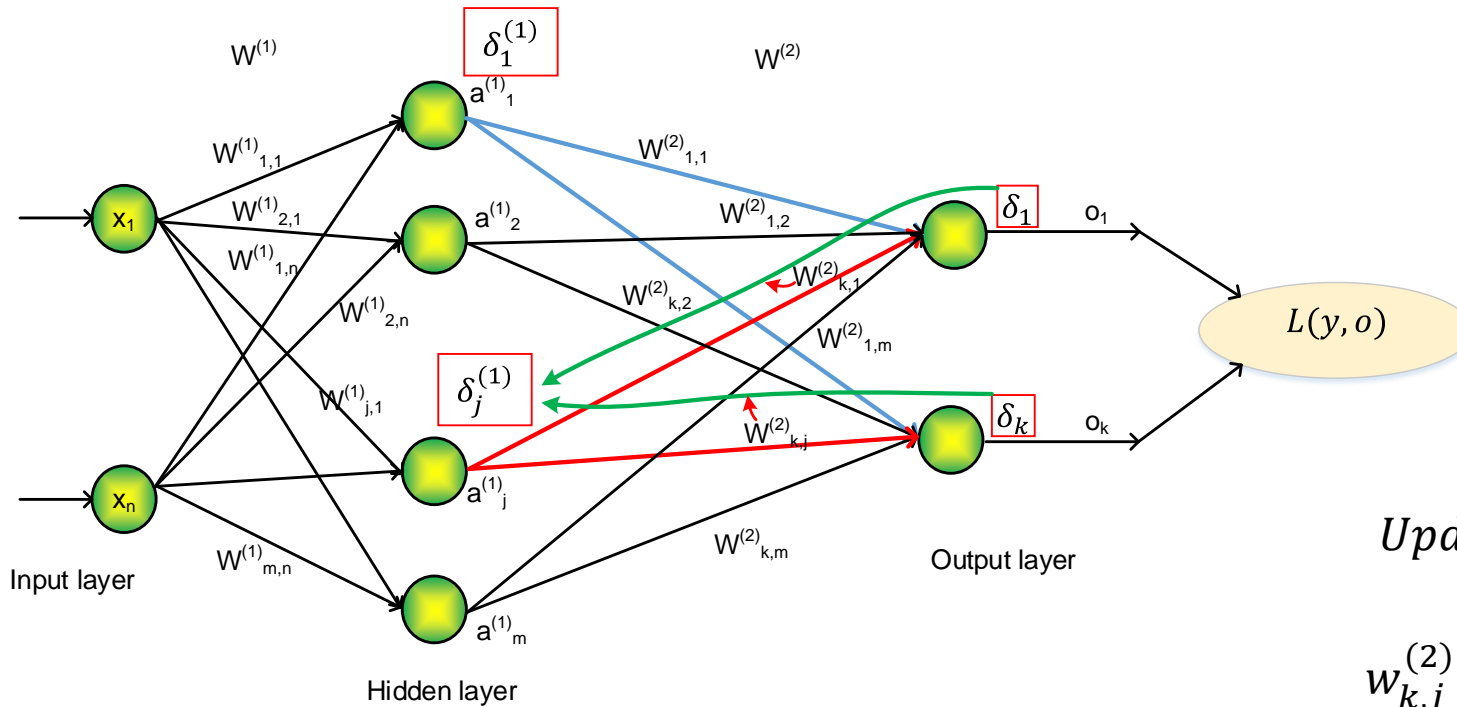
$\delta_j^{(1)}$ is the error of the i th neuron in hidden layer



$$\frac{\partial L}{\partial a_j^{(1)}} = \delta_j^{(1)}$$

$\delta_j^{(1)}$ affects all neurons of next layer via the weights

Back-propagation



Back propagate error

$$\delta_1^{(1)} = \delta_1 w_{1,1}^{(2)} + \dots + \delta_k w_{k,1}^{(2)}$$

$$\delta_j^{(1)} = \sum_{k=1}^K \delta_k w_{k,j}^{(2)}$$

Back propagate through sigmoid function

$$\frac{\partial L}{\partial z_j^{(1)}} = \delta_j^{(1)} a_j^{(1)} (1 - a_j^{(1)})$$

Update weights $w_{j,i} = w_{j,i} - \eta \frac{\partial L}{\partial w_{j,i}}$

$$w_{k,j}^{(2)} = w_{k,j}^{(2)} - \eta \frac{\partial L}{\partial w_{k,j}^{(2)}} = w_{k,j}^{(2)} - \eta \delta_k \frac{\partial \delta_k}{\partial w_{k,j}^{(2)}}$$

$$w_{k,j}^{(2)} = w_{k,j}^{(2)} - \eta \delta_k a_j^{(1)}$$

$$\delta_k = o_k - y_k$$

$$\delta_j^{(1)} = \sum_{k=1}^K \delta_k w_{k,j}^{(2)}$$

$$w_{j,i}^{(1)} = w_{j,i}^{(1)} - \eta \delta_j^{(1)} a_j^{(1)} (1 - a_j^{(1)}) x_i$$




you should spend time to master them


Stochastic gradient descent – batch and mini-batch gradient descent

- ❑ Stochastic gradient descent (SGD): use 1 sample in each iteration
- ❑ Batch gradient descent (GD): use all samples in each iteration
- ❑ Mini-batch gradient descent (Mini-batch GD): use b sample in each iteration, in this case, b is the batch size
- ❑ Mini-batch stochastic gradient descent (Mini-batch SGD): use b sample in each iteration, the batch of training samples is randomly selected

Stochastic Gradient Descent



Mini-Batch Gradient Descent




PYTHON CODE

Translating Math into Code




Translating mathematics into code

MNIST Dataset



Feed forward propagation



$$z2 = z^1$$

$$Wh = W^{(1)}$$

$$Wo = W^{(2)}$$

Define:

$$z2 = z^1$$

$$Wh = W^{(1)}$$

$$Wo = W^{(2)}$$

Forward

$$a = XWh^T$$

$$z1 = \sigma(z1) = \frac{1}{1 + e^{-z1}}$$

$$z2 = XWo^T$$

$$o_k = \frac{e^{z2_k}}{\sum_{j=1}^K e^{z2_k}}$$

Back-propagation error

Parameters

- n: number of inputs
- m: number of neurons in hidden layer
- k: number of neurons in output layer
- t: number of training sample (batch_size)


(element-wise product)

$$d_j^t = (o_j^t - y_j^t)$$


- ▶ Training sample # 1
- ▶ The k^{th} output



Outputs of neurons
number of training sample



$$d_j^t = (o_j^t - y_j^t)$$



$\times =$

$$W^{(2)}$$

$W_{1,1}$	$W_{1,2}$	$W_{1,3}$	$W_{1,m}$
$W_{2,1}$	$W_{2,2}$	$W_{2,3}$	$W_{2,m}$
$W_{k,1}$	$W_{k,2}$	$W_{k,3}$	$W_{k,m}$

 txk

күн


$$dh = [\delta_1^{(1)}, \delta_2^{(1)}, \dots, \delta_j^{(1)}, \dots, \delta_m^{(1)}]$$

$$\delta_j^{(1)} = \sum_{k=1}^K \delta_k w_{k,j}^{(2)}$$

$$dh = dW^{(2)}$$



txm



▶ Hidden layer

$$dh = [\delta_1^{(1)}, \delta_2^{(1)}, \dots, \delta_j^{(1)}, \dots, \delta_m^{(1)}]$$

 $d =$

Output layer




Diagram illustrating the structure of a weight matrix $W^{(1)}$ of size $m \times n$. The matrix is shown as a table with m rows and n columns. The first row is labeled n_1 and the last row is labeled n_m . The first column is labeled 1 and the last column is labeled n . The elements are $W_{1,1}, W_{1,2}, W_{1,3}, \dots, W_{1,n}$ in the first row, and $W_{m,1}, W_{m,2}, W_{m,3}, \dots, W_{m,n}$ in the last row.

Diagram illustrating the structure of the weight matrix $W^{(2)}$ of size $k \times m$. The matrix is shown as a table with k rows and m columns. The rows are indexed by n_1, \dots, n_k , and the columns are indexed by $1, 2, 3, \dots, m$. A red arrow points to the first column, indicating the input vector x .

n_1	$W_{1,1}$	$W_{1,2}$	$W_{1,3}$	$W_{1,m}$
\vdots	\vdots	\vdots	\vdots	\vdots
n_k	$W_{k,1}$	$W_{k,2}$	$W_{k,3}$	$W_{k,m}$

```
Wh = np.matrix(np.random.uniform
(-0.5,0.5,(NumHiddenUnits,NumInputs)))
```

```
Wo = np.random.uniform(
    -0.5,0.5,(NumClasses,NumHiddenUnits))
```

Back propagate error

$$\delta_1^{(1)} = \delta_1 w_{1,1}^{(2)} + \dots + \delta_k w_{k,1}^{(2)}$$

$$\delta_j^{(1)} = \sum_{k=1}^K \delta_k w_{k,j}^{(2)}$$

Update weights

output layer

$$L(y, o) = - \sum_{k=1}^K y_k \log(o_k) \quad \text{loss.append}(-\text{np.sum}(\text{np.multiply}(y, \text{np.log10}(o))))$$

$$d = o - y \quad \text{softmax} \quad d = o - y$$

$$\Delta W_o = -\eta \frac{2}{t} d^T a \quad dWo = \text{np.matmul}(\text{np.transpose}(d), a)$$

$$W_o = W_o + \Delta W_o$$

output layer

$$dh = dW_o \quad dh = d @ W_o \quad \text{back propagate through weights}$$

$$dhs = dh(a(1-a)) \quad \text{back propagate through sigmoid}$$

$$dhs = \text{np.multiply}(\text{np.multiply}(dh, a), (1-a))$$


$$\Delta W_o = -\eta \frac{2}{t} dhs^T X$$

t : number of training sample

Update weights

$$w_{j,i} = w_{j,i} - \eta \frac{\partial L}{\partial w_{j,i}}$$

$$w_{k,j}^{(2)} = w_{k,j}^{(2)} - \eta \frac{\partial L}{\partial w_{k,j}^{(2)}} = w_{k,j}^{(2)} - \eta \delta_k \frac{\partial \delta_k}{\partial w_{k,j}^{(2)}}$$



PYTHON CODE



Python code

Load dataset

Batch Gradient Descent

```
import numpy as np
import tensorflow as tf
#load dataset
print("Load MNIST Database")
mnist = tf.keras.datasets.mnist
(x_train,y_train),(x_test,y_test)= mnist.load_data()
x_train=np.reshape(x_train,(60000,784))/255.0
x_test= np.reshape(x_test,(10000,784))/255.0
y_train = np.matrix(np.eye(10)[y_train])
y_test = np.matrix(np.eye(10)[y_test])
print("-----")
print(x_train.shape)
print(y_train.shape)
```

Python code

Define functions

```
def sigmoid(x):  
    return 1./(1.+np.exp(-x))  
  
def softmax(x):  
    return np.divide(np.matrix(np.exp(x)),np.mat(np.sum(np.exp(x),axis=1)))  
  
def Forwardpass(X,Wh,bh,Wo,bo):  
    zh = X@Wh.T + bh  
    a = sigmoid(zh)  
    z=a@Wo.T + bo  
    o = softmax(z)  
    return o  
  
def AccTest(label,prediction):    # calculate the matching score  
    OutMaxArg=np.argmax(prediction,axis=1)  
    LabelMaxArg=np.argmax(label,axis=1)  
    Accuracy=np.mean(OutMaxArg==LabelMaxArg)  
    return Accuracy
```

Python code

Define network architecture, initialize weights

```
learningRate = 0.5
Epoch=50
NumTrainSamples=60000
NumTestSamples=10000

NumInputs=784
NumHiddenUnits=512
NumClasses=10
#inital weights
#hidden layer
Wh=np.matrix(np.random.uniform(-0.5,0.5,(NumHiddenUnits,NumInputs)))
bh= np.random.uniform(0,0.5,(1,NumHiddenUnits))
dWh= np.zeros((NumHiddenUnits,NumInputs))
dbh= np.zeros((1,NumHiddenUnits))
#Output layer
Wo=np.random.uniform(-0.5,0.5,(NumClasses,NumHiddenUnits))
bo= np.random.uniform(0,0.5,(1,NumClasses))
dWo= np.zeros((NumClasses,NumHiddenUnits))
dbo= np.zeros((1,NumClasses))
```



Python code

Batch Gradient Descent

Training the model

```
from IPython.display import clear_output
loss = []
Acc = []
for ep in range (Epoch):
    #feed forward propagation
    x = x_train
    y=y_train
    zh = x@Wh.T + bh
    a = sigmoid(zh)
    z=a@Wo.T + bo
    o = softmax(z)
    #calculate loss
    loss.append(-np.sum(np.multiply(y,np.log10(o))))
    #calculate the error for the ouput layer
    d = o-y
    #Back propagate error
    dh = d@Wo
    dhs = np.multiply(np.multiply(dh,a),(1-a))
    #update weight
    dWo = np.matmul(np.transpose(d),a)
    dbo = np.mean(d) # consider a is 1 for bias
    dWh = np.matmul(np.transpose(dhs),x)
    dbh = np.mean(dhs) # consider a is 1 for bias
    Wo =Wo - learningRate*dWo/NumTrainSamples
    bo =bo - learningRate*dbo
    Wh =Wh-learningRate*dWh/NumTrainSamples
    bh =bh-learningRate*dbh
    #Test accuracy with random innitial weights
    prediction = Forwardpass(x_test,Wh,bh,Wo,bo)
    Acc.append(AccTest(y_test,prediction))
    clear_output(wait=True)
    plt.plot([i for i, _ in enumerate(Acc)],Acc,'o')
    plt.show()
```

Python code



Test the model

```
prediction = Forwardpass(x_test,Wh,bh,Wo,bo)
Rate = AccTest(y_test,prediction)
Print(Rate)
```

Python code

Training the model

```
from IPython.display import clear_output
loss = []
Acc = []
Batch_size = 200
Stochastic_samples = np.arange(NumTrainSamples)
for ep in range (Epoch):
    np.random.shuffle(Stochastic_samples)
    for ite in range (0,NumTrainSamples,Batch_size):
        #feed forward propagation
        Batch_samples = Stochastic_samples[ite:ite+Batch_size]
        x = x_train[Batch_samples,:]
        y=y_train[Batch_samples,:]
        zh = x@Wh.T + bh
        a = sigmoid(zh)
        z=a@Wo.T + bo
        o = softmax(z)
        #calculate loss
        loss.append(-np.sum(np.multiply(y,np.log10(o))))
        #calculate the error for the ouput layer
        d = o-y
        #Back propagate error
        dh = d@Wo
        dhs = np.multiply(np.multiply(dh,a),(1-a))
        #update weight
```

Mini-Batch Gradient Descent

```
#update weight
dWo = np.matmul(np.transpose(d),a)
dbo = np.mean(d) # consider a is 1 for bias
dWh = np.matmul(np.transpose(dhs),x)
dbh = np.mean(dhs) # consider a is 1 for bias
Wo =Wo - learningRate*dWo/Batch_size
bo =bo - learningRate*dbo
Wh =Wh-learningRate*dWh/Batch_size
bh =bh-learningRate*dbh
#Test accuracy with random innitial weights
prediction = Forwardpass(x_test,Wh,bh,Wo,bo)
Acc.append(AccTest(y_test,prediction))
clear_output(wait=True)
plt.plot([i for i, _ in enumerate(Acc)],Acc,'o')
plt.show()
print('Epoch:', ep )
print('Accuracy:',AccTest(y_test,prediction) )
```

Python code

Just calculate the loss and accuracy after each epoch

Mini-Batch Gradient Descent

```
from IPython.display import clear_output
loss = []
Acc = []
Batch_size = 200
Stochastic_samples = np.arange(NumTrainSamples)
for ep in range (Epoch):
    np.random.shuffle(Stochastic_samples)
    for ite in range (0,NumTrainSamples,Batch_size):
        #feed forward propagation
        Batch_samples =
Stochastic_samples[ite:ite+Batch_size]
        x = x_train[Batch_samples,:]
        y=y_train[Batch_samples,:]
        zh = x@Wh.T + bh
        a = sigmoid(zh)
        z=a@Wo.T + bo
        o = softmax(z)
        #calculate loss
        loss.append(-np.sum(np.multiply(y,np.log10(o))))
        #calculate the error for the ouput layer
        d = o-y
        #Back propagate error
        dh = d@Wo
        dhs = np.multiply(np.multiply(dh,a),(1-a))
```

```
dWo = np.matmul(np.transpose(d),a)
dbo = np.mean(d) # consider a is 1 for bias
dWh = np.matmul(np.transpose(dhs),x)
dbh = np.mean(dhs) # consider a is 1 for bias
Wo =Wo - learningRate*dWo/Batch_size
bo =bo - learningRate*dbo
Wh =Wh-learningRate*dWh/Batch_size
bh =bh-learningRate*dbh
#Test accuracy with random innitial weights
prediction = Forwardpass(x_test,Wh,bh,Wo,bo)
Acc.append(AccTest(y_test,prediction))
print('Epoch:', ep )
print('Accuracy:',AccTest(y_test,prediction) )
```

```
Epoch: 0
Accuracy: 0.8762
Epoch: 1
Accuracy: 0.9013
Epoch: 2
Accuracy: 0.9136
Epoch: 3
Accuracy: 0.9165
Epoch: 4
Accuracy: 0.9251
```



Python code

Training the model

Mini-Batch Gradient Descent

