

AI - FOUNDATION AND APPLICATION

Instructor:

Assoc. Prof. Dr. Truong Ngoc Son

Chapter 4

Convolution Neural Network

Outline



Motivation – Convolution operation

Convolution is a function derived from two given functions by integration which expresses how the shape of one is modified by the other

$$y(t) = x(t)(*)h(t) = \int_{\tau=-\infty}^{+\infty} x(\tau)h(t-\tau)d\tau$$

Discrete convolution

$$y(t) = x(t)(*)h(t) = \sum_{\tau=-\infty}^{+\infty} x(\tau)h(t-\tau)$$


Motivation – Convolution operation

Convolution in 2-D space


$$y(n_1, n_2) = \sum_{k_2=0}^{N-1} \sum_{k_1=0}^{M-1} x(k_1, k_2) h(n_1 - k_1, n_2 - k_2)$$

$$0 \leq n_1 \leq N - 1, 0 \leq n_2 \leq M - 1.$$

Computer Vision Problem



vertical edges



horizontal edges

Vertical edge detection

3 ¹	0 ⁰	1 ⁻¹	2 ⁻¹	7 ⁻⁰	4 ⁻¹
1 ¹	5 ⁰	8 ⁻¹	9 ⁻¹	3 ⁻⁰	1 ⁻¹
2 ¹	7 ⁰	2 ⁻¹	5 ⁻¹	1 ⁻⁰	3 ⁻¹
0 ¹	1 ⁰	3 ⁻¹	1 ⁻¹	7 ⁻⁰	8 ⁻¹
4	2	1	6	2	8
2	4	5	2	3	9

*

1	0	-1
1	0	-1
1	0	-1

=

-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

Vertical edge detection


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*



1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



*



Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



*

1	0	-1
1	0	-1
1	0	-1



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10



*

1	0	-1
1	0	-1
1	0	-1




=


0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



Convolution Neural Network




(a) Fully connector layer




(b) Convolution layer

Convolution Neural Network






Convolution Neural Network


- Input layer
- Convolution layer
- Activation function
- Pooling layer
- Fully connected layers




Convolution layer




Padding



*





Valid and Same convolutions

“Valid”: No padding is required

“Same”: Pad so that output size is the same as the input size.

Strided convolution²

2 ³	3 ⁴	7 ³	4 ⁴	6 ³	2 ⁴	9 ⁴
6 ¹	6 ⁰	9 ¹	8 ⁰	7 ¹	4 ⁰	3 ²
3 ⁻³	4 ⁴	8 ³	3 ⁴	8 ³	9 ⁴	7 ⁴
7 ¹	8 ⁰	3 ¹	6 ⁰	6 ¹	3 ⁰	4 ²
4 ⁻³	2 ⁴	1 ³	8 ⁴	3 ³	4 ⁴	6 ⁴
3 ¹	2 ⁰	4 ¹	1 ⁰	9 ¹	8 ⁰	3 ²
0 ⁻¹	1 ⁰	3 ⁻¹	9 ⁰	2 ⁻¹	1 ⁰	4 ³

*

3	4	4
1	0	2
-1	0	3

=

Summary of convolutions

$n \times n$ image $f \times f$ filter

padding p stride s

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

Technical note on cross-correlation vs. convolution

Convolution in math textbook:

2	3	7	4	6	2
6	6	9	8	7	4
3	4	8	3	8	9
7	8	3	6	6	3
4	2	1	8	3	4
3	2	4	1	9	8

*

3	4	5
1	0	2
-1	9	7

7	2	5
9	0	4
-1	1	3

Pooling layer: Max pooling

1	3	2	1
2	9	1	1
1	3	2	3
5	6	1	2

9	2
6	3

Pooling layer: Max pooling

1	3	2	1	3
2	9	1	1	5
1	8	2	3	2
8	3	5	1	0
5	6	1	2	9

Pooling layer: Average pooling

1	3	2	1
2	9	1	1
1	4	2	3
5	6	1	2




3.75	1.25
4	2

Fully-connected layer (FC)

Flatten

1	1	0
4	2	1
0	2	1



1
1
0
4
2
1
0
2
1

Summary of pooling

Hyperparameters:

f : filter size

s : stride

Max or average pooling

CNN

- 1. Lenet-5
- 2. Alex-net
- 3. VGG
- 4. ResNet
- 5. Inception
- 6. MobileNet
- 7. Data Augmentation
- 8. Object detection: R-CNN
- 9. Object detection: Yolo
- Transfer Learning ()

PYTHON CODE

